

SAIGON BUSINESS SCHOOL

Department of Computer Science



**Saigon
Business
School**

FINAL ASSIGNMENT

Full Name: Phone Myat Min

Student ID: SBS24010097

Section class: Data Structures & Algorithms

Instructor: MSc. Luong Tran Ngoc Khiet

Submission date: 09 May 2026

HO CHI MINH CITY, 2026

CHECKLIST

Please check all of the following items before submitting your work. If any of the following items are missing or incorrect, your work may not be accepted

No	Checklist	Yes or No
1	I have checked the accuracy of the information on the cover page (course name, student name, etc.)	Yes
2	I have read and signed the Statement of ownership and integrity.	Yes
3	I have created a table of contents for my work.	Yes
4	I have included my name in the footer.	Yes
5	I have used the correct font and font size as instructed for the body of my work*.	Yes
7	I have cited sources in my work using the Harvard Reference format.	Yes
8	I have created a reference list in the correct format.	Yes

* Format requirements

- Use Times New Roman font, size 13, line spacing 1.5;
- Margins: Left 3cm, Right 2cm, Top 2cm, Bottom 2cm;
- All references must be properly cited and listed in the References section.

PART 1 –ASSESSMENT SUMMARY

Learning outcome	Criteria	The student provides evidence for their work:	Page	Result (Yes/No)
<u>Learning Outcome 1:</u>	P1	Create a design specification for data structures, explaining the valid operations that can be carried out on the structures	8-12	
<u>Learning Outcome 1:</u>	P2	Determine the operations of a memory stack and how it is used to implement function calls in a computer.	21-23	
<u>Learning Outcome 1:</u>	M1	Illustrate, with an example, a concrete data structure for a First in First out (FIFO) queue.	12-14	
<u>Learning Outcome 1:</u>	M2	Compare the performance of two sorting algorithms.	14-18	
<u>Learning Outcome 1:</u>	D1	Analyze the operation, using illustrations, of two network shortest path algorithm, providing a example of each.	25-27	
<u>Learning Outcome 2:</u>	P3	Specify the abstract data type for a software stack using an imperative definition.	23-25	
<u>Learning Outcome 2:</u>	M3	Examine the advantages of encapsulation and information hiding when using an ADT.	18-19	
<u>Learning Outcome 2:</u>	D2	Discuss the view that imperative ADTs are a basis for object orientation offering a justification for the view.	19-20	
<u>Learning Outcome 3:</u>	P4	Implement a complex ADT and algorithm in a executable programming language to solve a well-defined problem.	31-35	

Learning outcome	Criteria	The student provides evidence for their work:	Page	Result (Yes/No)
<u>Learning Outcome 3:</u>	P5	Implement error handling and report test results.	35-37	
<u>Learning Outcome 3:</u>	M4	Demonstrate how the implementation of an ADT/algorithm solves a well-defined problem.	31-35	
<u>Learning Outcome 3:</u>	D3	Critically evaluate the complexity of an implemented ADT/algorithm.	37-39	
<u>Learning Outcome 4:</u>	P6	Discuss how asymptotic analysis can be used to assess the effectiveness of an algorithm.	27-28	
<u>Learning Outcome 4:</u>	P7	Determine two ways in which the efficiency of an algorithm can be measured, illustrating your answer with an example.	28-30	
<u>Learning Outcome 4:</u>	M5	Interpret what a trade-off is when specifying an ADT, using an example to support your answer.	30-31	
<u>Learning Outcome 4:</u>	D4	Evaluate three benefits of using implementation independent data structures.	39-41	

PART 2 – LECTURER’S FEEDBACK

Formative Feedback:
Action Plan

Student Feedback to Assessor			
Summative Feedback			
Student's signature	Phone Myat Min	Date	09 May 2026
Assessor's signature		Date	

Contents

Task 1: System Architecture and Cargo Sorting.....	8
P1: Design Specification for the "Goods" Data Structure.....	8
M1: Queue Illustration (FIFO) for Loading Bay Management.....	12
M2: Sorting Comparison (Bubble Sort vs QuickSort)	14
M3: Data Encapsulation and Information Hiding.....	18
D2: OOP Foundations from Imperative ADTs	19
Task 2: Operational Optimization and Efficiency.....	20
P2: Stack Operations for Process Control.....	21
P3: Formal Specification for a Software Stack (Imperative Definition)	23
D1: Routing Analysis – Dijkstra’s vs A* Algorithm.....	25
P6: Effectiveness Assessment via Big O Notation.....	27
P7: Two Methods for Measuring Algorithm Efficiency	28
M5: Trade-off Interpretation in ADT Specification.....	30
Task 3: Implementation and Technical Critique.....	31
P4 & M4: Complex ADT Implementation – AVL Tree for Warehouse Inventory.....	31
P5: Robustness and Testing Report	35

D3: Critical Evaluation – AVL Tree vs Linear Search Using Big O	37
D4: Three Benefits of Implementation-Independent Data Structures.....	39
References.....	41
Appendix.....	41

Task 1: System Architecture and Cargo Sorting

This task is about designing the first version of Swift-Load Logistics platform: a tool that will manage cargo items in a highly efficient way. We create the main Abstract Data Type (ADT) for goods, demonstrate truck management with the help of queue, compare sorting algorithms for cargo optimization, and discuss how principles of encapsulation keep the system safe from outside interference.

P1: Design Specification for the "Goods" Data Structure

The Goods Abstract Data Type is the backbone of the Swift-Load Logistics system. In general, an ADT specifies a data type mathematically through 2 main points: the set of values it can take, and the set of operations that can be done on it. To put it simply, an ADT is a triple (D F A) where D is the domain of values, F is the set of functions/operations, and A is the set of axioms/constraints that determine the behavior.

Attributes (Domain D):

- **Name (String):** Ordinarily, each item is recognized by a unique, human-readable name, e.g. "Laptop" or "Apples". It also acts as the key when the items are stored in a tree structure and is the main criterion for alphabetical sorting and comparison through binary search. The name must be non-empty to be inserted correctly into the AVL tree.
- **Type (String):** Item type (e.g. "Electronics", "Food", "Furniture"). This will allow grouping of cargo, recognition of special handling requirements such as refrigeration for food items, and also priority-based loading sequences where, for example, fragile electronics are loaded last.

- **Weight (Double):** Item weight in kilograms. This information is crucial for truck capacity planning, load balancing calculations, and determining whether a truck has been overloaded beyond its maximum weight capacity. The weight should always be positive since a negative weight does not physically exist and would ruin further calculations.

Valid Operations (Functions F):

The below table outlines all the valid operations that a Goods ADT can perform. Each one comes with its parameter and return types. Besides, validation constraints that guarantee data consistency in the system are included. These operations are the entire public interface by which Goods objects are accessed by all outside modules.

Operation	Parameters	Return Type	Pre/Post Conditions & Axioms
Constructor (init)	name: String, type: String, weight: Double	N/A	Post: Valid Goods object created in memory. Axiom: Exceptions are evaluated in order: Name first, then Weight. If name is empty, throws <code>invalid_argument</code> . If <code>weight < 0</code> , throws <code>invalid_argument</code> .
<code>getName()</code>	None	String	Post: Returns the current name value without modifying internal state (const method).
<code>setName(newName)</code>	<code>newName: String</code>	void	Pre: <code>newName</code> must not be empty string. Post: Internal name updated to <code>newName</code> . Axiom: if <code>newName</code> is empty, throw

			invalid_argument exception.
getType()	None	String	Post: Returns the current type classification string.
setType(newType)	newType: String	void	Post: Internal type updated to newType. No restrictions on value.
getWeight()	None	Double	Post: Returns the current weight value in kilograms as a double-precision float.
setWeight(newWeight)	newWeight: Double	void	Pre: newWeight >= 0. Post: Internal weight updated. Axiom: if newWeight < 0, throw invalid_argument exception to prevent data corruption.
display()	None	void	Post: Outputs formatted Name (20 chars), Type (15 chars), Weight to console with "kg" suffix.

Table 1: Goods ADT Complete Operation Specification

In the C++ code, the getter methods (such as getName(), getWeight()) receive the const specifier, and this is actually a quite significant software engineering practice. It not only assures someone's ADT can be passed by reference to other functions (like sorting or searching algorithms) with no risk of mutation, but also the compiler will make sure that no changes can be made to the object through the const reference, This way preventing bugs in the logistics pipeline that are hard to trace.

C++ Implementation of the Goods Class:

```
class Goods {
private:
    string name;
    string type;
    double weight;

public:
    Goods(string name = "Unnamed", string type = "", double weight = 0.0) {
        // Validation order: Name first, then Weight
        if (name.empty()) {
            throw invalid_argument("Name cannot be empty");
        }
        if (weight < 0) {
            throw invalid_argument("Weight cannot be negative");
        }
        this->name = name;
        this->type = type;
        this->weight = weight;
    }

    // --- Accessor Methods (Getters) ---
    string getName() const { return name; }
    string getType() const { return type; }
    double getWeight() const { return weight; }

    // --- Mutator Methods (Setters) with Validation ---
    void setName(const string& newName) {
        if (newName.empty())
            throw invalid_argument("Name cannot be empty");
        name = newName;
    }
    void setType(const string& newType) { type = newType; }
    void setWeight(double newWeight) {
        if (newWeight < 0) {
            throw invalid_argument("Weight cannot be negative");
        }
    }
}
```

```

    }
    weight = newWeight;
}
};

```

The constructor shown above is a good example of how one can enforce the precondition that weight has to be a non-negative number. When a negative weight is passed to it, a C++ `std::invalid_argument` exception is thrown so corrupted data are prevented from getting into the logistics pipeline. Also, a `const` qualifier on getter methods is a promise that a call to `getName()` shouldn't change the object's internal state. This acts as another layer of safety that compilers can check. Such a defensive approach to programming is crucial for the integrity of the data throughout the whole Swift-Load platform, especially when the same Goods instances are accessed concurrently by multiple modules.

M1: Queue Illustration (FIFO) for Loading Bay Management

To keep things running smoothly and fairly, the Swift-Load loading bay gives priority to trucks on a strict first-come-first-served basis. The whole process is mathematically modeled as a First-In-First-Out (FIFO) Queue. In essence, a queue is a linear data structure that allows two main operations: enqueue which adds an element to the rear and dequeue which removes an element from the front. Unlike a stack that accesses the most recently added element first, a queue will always process elements in the exact order in which they arrived, thus preventing a situation where a truck could be left waiting indefinitely while newer arrivals are processed first.

Consider, in a logistics scenario, if Truck A gets to the loading bay at 8:00 AM, followed by Truck B at 8:15 AM, and Truck C at 8:30 AM, using a queue will enforce Truck A always being the one that loads first. Moreover, keeping to such FIFO discipline is extremely important for ensuring delivery deadlines that have been contracted are met as well as for customer satisfaction purposes.

Concrete Step-by-Step Example:

Step	Operation	Queue State (Front - > Rear)	Queue Size
1	Enqueue Truck A (Heavy Duty)	[Truck A]	1
2	Enqueue Truck B (Refrigerated)	[Truck A, Truck B]	2
3	Enqueue Truck C (Standard)	[Truck A, Truck B, Truck C]	3
4	Dequeue -> Truck A processed	[Truck B, Truck C]	2
5	Dequeue -> Truck B processed	[Truck C]	1
6	Dequeue -> Truck C processed	[] (empty)	0

Table 2: Step-by-Step FIFO Queue Operations for Loading Bay

Queue Diagram (Visual Representation):

EXIT <-- Dequeue	Truck A (Front)	Truck B (Middle)	Truck C (Rear)	ENTER Enqueue -->
---------------------	--------------------	---------------------	-------------------	----------------------

Figure 1: FIFO Queue Diagram for Swift-Load Loading Bay

Though it is a common approach, using a FIFO queue in logistics is not a very good idea from the point of view of operations, since it ignores the priority of the cargo. For instance, if the life-saving perishable medical supplies carried by Truck C were to arrive, while Trucks A and B had earlier brought non-urgent scrap metal, then as the standard FIFO the medical supplies would have to wait. Whilst FIFO queue basically meets the chronological requirement as given in the assignment, a Priority Queue (implemented via a Max-Heap or Min-Heap) would be a huge step forward for Swift-Load. Such a priority queue might allow the system to work on trafficking high-priority items right after their arrival, irrespective of the length of the queue. It is necessary to be aware of the

shortcomings of standard linear chronologies when one is doing enterprise architectural design.

C++ Implementation using std::queue:

```
queue<string> truckQueue;
truckQueue.push("Truck A (Heavy Duty)");    // Enqueue at rear
truckQueue.push("Truck B (Refrigerated)");  // Enqueue at rear
truckQueue.push("Truck C (Standard)");      // Enqueue at rear
// Front: Truck A, Back: Truck C, Size: 3
while (!truckQueue.empty()) {
    cout << "Processing: " << truckQueue.front(); // Peek at front
    truckQueue.pop();                             // Dequeue from front
}
```

Among the data structure classes supplied by the C++ Standard Template Library is the `std::queue` container adapter that realizes the FIFO concept. Thus, the `push()` method enqueues, `front()` peeks, and `pop()` dequeues. Besides, as the time complexity of all these three functions is $O(1)$, which is constant, the queue becomes very efficient even when managing a loading bay in real-time in case of many trucks waiting.

M2: Sorting Comparison (Bubble Sort vs QuickSort)

Efficiently utilizing a truck's capacity and ensuring that the distribution of the load is safe require the sorting of cargo by weight. First, heavier items should be loaded (go to the bottom) in order to ascertain safety and make the best use of space. We took twelve randomly selected cargo weights to illustrate loading a heavy-first approach and thus, as a working example, implemented and compared two completely different sorting algorithms.

Sample Data (12 cargo weights in kg):

{45.5, 12.0, 89.2, 5.5, 34.1, 100.0, 77.8, 23.4, 56.7, 9.9, 41.2, 60.0}

Expected Sorted Result (Descending):

{100.0, 89.2, 77.8, 60.0, 56.7, 45.5, 41.2, 34.1, 23.4, 12.0, 9.9, 5.5}

Bubble Sort Algorithm:

Bubble sort, a simple algorithm based on the principle of comparison, continuously passes through the list, compares each pair of adjacent elements and if they are found to be in the wrong order, simply exchanges their places. At the end of each pass, the smallest element among those remaining unsorted will "have bubbled" up to its correct position (as a bubble rises to the surface). It is clear that in the worst-case scenario $N-1$ passes will be required. Each pass will be performing up to $N-i-1$ comparisons. Although it is the easiest to understand and the implementation straightforward, because of its time complexity being $O(N^2)$, it is normally considered impractical for input sizes exceeding a few hundred items.

```
void bubbleSort(vector<double>& arr, long long& comps, long long& swaps) {
    size_t n = arr.size();
    for (int i = 0; i < n - 1; i++) {
        for (int j = 0; j < n - i - 1; j++) {
            if (arr[j] < arr[j + 1]) { // < for descending
                swap(arr[j], arr[j + 1]);
            }
        }
    }
}
```

QuickSort Algorithm:

QuickSort employs a divide-and-conquer approach. It chooses a pivot element (usually the last one), divides the array into parts containing elements are bigger than and smaller than the pivot (for descending order), and then sorts each part separately through recursion. QuickSort's distinguishing feature is partitioning: it changes the positions of the elements with a single linear pass, and then the two sub-problems generated are more or less half the original size. This behavior of halving the work leads to QuickSort's $O(N \log$

N) average-case performance, which makes it significantly faster than Bubble Sort when working with large datasets.

```
void quickSort(vector<double>& arr, int low, int high) {
    if (low < high) {
        int pi = partition(arr, low, high); // Partition
        quickSort(arr, low, pi - 1);      // Sort left half
        quickSort(arr, pi + 1, high);     // Sort right half
    }
}
```

Empirical Performance Comparison:

Metric	Bubble Sort	QuickSort	Winner	Significance for Logistics
Time Complexity (Avg)	$O(N^2)$	$O(N \log N)$	QuickSort	For N=10,000 packages: Bubble=100M ops vs Quick=133K ops
Time Complexity (Worst)	$O(N^2)$	$O(N^2)$	Tie	QuickSort worst case occurs with already-sorted input
Comparisons (12 items)	66	34	QuickSort	QuickSort uses 48% fewer comparisons on our sample
Swaps (12 items)	34	23	QuickSort	QuickSort performs 32% fewer memory write operations
Space Complexity	$O(1)$ in-place	$O(\log N)$ stack	Bubble Sort	QuickSort requires recursive stack frames

Stability	Stable	Unstable	If two packages weigh exactly 50.0kg, a stable sort (Bubble) preserves the order they arrived in the system. An unstable sort (Quick) might arbitrarily swap them. If arrival time dictates loading priority for identical weights, QuickSort's instability is actually a minor drawback that must be accounted for.	Bubble Sort preserves relative order of equal elements
Implementation	Very Simple	Moderate	Bubble Sort	Bubble Sort has only 2 nested loops, easier to debug

Table 3: Comprehensive Bubble Sort vs QuickSort Performance Comparison

Summary: Even though Bubble Sort has some slight advantages in terms of space efficiency and stability, QuickSort is, by far, the better choice for the Swift-Load scenario. Since thousands of packages are handled daily, the $O(N^2)$ time complexity of Bubble Sort would cause unacceptable delays during the high-season times, whereas QuickSort's $O(N \log N)$ keeps the input/output flow steady no matter the amount of goods. According to our tests, QuickSort makes 48% fewer comparisons even with only 12 items.

M3: Data Encapsulation and Information Hiding

Encapsulation is among the four main principles of Object-Oriented Programming, together with inheritance, polymorphism, and abstraction. It is the concept of packaging data (attributes) and functions that operate on the data together into one unit (class), while also limiting the direct access to certain components of the object. In the Goods abstract data type (ADT), all three attributes (name type weight) are made private in C++, so they can neither be accessed nor altered directly from outside the class. All the work has to be done via the public methods (getters and setters) that have been checked and validated. This is an essential measure for maintaining security and integrity in the logistics system.

Advantage 1: Data Integrity Through Validation Gates

With no encapsulation, any part of the code might run `cargo.weight = -500` directly, which would instantly break the load-balancing algorithm and, quite likely, lead to a truck being either overloaded or underloaded. Making all weight changes via the `setWeight()` method only, we ensure that the validation (if `newWeight < 0`, throw exception) is always carried out before the internal state actually changes. This way, there is a single validation checkpoint that safeguards the entire system against incorrect data, no matter which module makes the change.

Advantage 2: Implementation Independence and Future-Proofing

Later on, if the logistics department opts to keep weight in grams instead of kilograms to be more accurate on the lightweight electronics, only the inner implementation of `getWeight()` and `setWeight()` will have to be changed. All the other modules reading weight will still operate exactly the same way since their interaction is only through the public interface. This lowers the expense and danger of changing the system so significantly.

Advantage 3: Reduced Coupling and Simplified Maintenance

Encapsulation is a way of implementing software that reduces the dependence of one component on another, simply by hiding the inner details from the users. The truck-loading software is not concerned about what type of variable weight is - double, float, or integer - internally. Such a modular architecture also assures that any errors in the Goods class stay contained there and would not cause the routing or billing modules to become corrupted. This sealed environment therefore helps a lot in reducing the time taken for debugging as well as the overall cost of maintenance throughout the software lifecycle.

D2: OOP Foundations from Imperative ADTs

The transformation from imperative Abstract Data Types (ADTs) to Object-Oriented Programming (OOP) marks one of the biggest paradigm changes in the history of computer science. In the case of pure imperative languages such as C, data structures (structs) and the functions that operate on them are totally separate. For example, a C struct for Goods would merely hold the data elements, while non-member functions like `updateWeight(Goods* g, double w)` would work on it from outside. Typically the code is not able to restrict direct access to individual struct members (`g->weight = -500`), cannot package the validation logic with the data it is supposed to preserve, and does not have a formal linkage between the struct and its associated functions.

Nonetheless, the idea of imperative ADT gave the key philosophical insight that allowed to close this gap: some data and operations are semantically inseparable. You cannot update the weight of a cargo item coming from one operation and the validation of the new value being non-negative must always be carried out along with it. Similarly, making the name of an item empty may result in breaking the AVL tree search key invariant. This awareness of data and operations that are permissible, together constitute a single logical unit, has become the philosophical as well as mathematical basis of OOP.

With our Swift-Load system, this evolutionary leap can be directly seen in the codebase. The imperative ADT specification declared: A Goods value is made up of name,

type, and weight. Valid operations are init getName setName getWeight setWeight, and display. The C++ class implementation mirrors this formal specification exactly into language-level constructs: private member variables, as the encapsulated data, publicly accessible methods, as the allowable operations constructors to formalize the initialization, and exception handling, to enforce the axioms. The C++ class IS the ADT, implemented in executable code. This direct correspondence is evidence that Object-Oriented Programming is not a separate invention but the natural and inevitable formalization of imperative ADT principles at the programming language level, and it is compiler-enforced guarantees that replace what was previously only a design convention.

Task 2: Operational Optimization and Efficiency

Task focuses on the theoretical and analytical dimensions of resource management, process control, and routing optimization in the Swift-Load Logistics system. Function execution is managed via stack-based process control. Stack ADT is formally specified in imperative notation. We conduct analysis of different routing algorithms, and finally, through asymptotic analysis and empirical measurement, we determine which of these algorithms is most effective.

Encapsulation exists fundamentally to protect the class invariants (the undeniable rules of the data structure, such as 'weight must always be ≥ 0 ' or 'name cannot be empty'). By hiding the data and exposing only validated setters, we prevent any external code from violating these invariants.

To visually understand the profound evolution from imperative C to Object-Oriented C++, consider this comparison:

```
// Imperative C-Style (No Encapsulation)
struct Goods { string name; double weight; };
// Anyone can do: myGoods.weight = -500; (DANGER)

// OOP C++ Style (Encapsulated ADT)
```

```

class Goods {
private:
    double weight;
public:
    void setWeight(double w) {
        if(w < 0) throw error();
        weight = w;
    }
};
// myGoods.weight = -500; // COMPILER ERROR

```

P2: Stack Operations for Process Control

Memory stacks operate using LIFO (Last In First Out) as do queues with the exception that while the first item in the queue is always the first item out of the queue, the last item entered into the memory stack will be the very next item to leave the memory stack. This is just like how a stack of plates works. The only place you may add or remove a plate from is at the "top" of the stack.

There are many ways for computer programs to manage their function executions. One of those methods is through use of a Call Stack. The Call Stack has two main operations; they are push and pop. When a new function is called, a new stack frame is created for this function and is then placed onto the top of the Call Stack. Each stack frame contains four pieces of information about the current function being executed by the CPU. They are: (1) the local variables associated with each function, (2) all parameters passed to a function when it was called, (3) the return address to determine where execution of the program should resume once a function completes execution, and (4) any registers that were used by the CPU prior to calling a particular function that need to be restored so that execution of another function can occur. Once the execution of a function has completed, the top stack frame of the Call Stack is removed (pop operation), and all information in that stack frame is destroyed. The Call Stack now looks like it did before that function was called and execution resumes at the point where it left off.

Core Stack Operations:

Operation	Description	Time Complexity	Application in Swift-Load
push(item)	Adds item to the top of the stack	O(1)	When quickSort() calls partition(), a new frame is pushed
pop()	Removes and returns the top item	O(1)	When partition() returns, its frame is popped
top()/peek()	Returns the top item without removing	O(1)	Checking which function is currently executing
isEmpty()	Returns true if stack has no elements	O(1)	Checking if all functions have returned
size()	Returns the number of elements	O(1)	Monitoring recursion depth for stack overflow prevention

Table 4: Stack Operations with Logistics Applications

Call Stack Diagram During QuickSort Execution:

Direction	PUSH ↓ down POP ↑ up
TOP →	partition(arr, 0, 5) frame [local vars: pivot, i, j]
	quickSort(arr, 0, 5) frame [local vars: pi]
	loadTruck(truckId=42) frame [local vars: items, capacity]
BOTTOM →	main() frame [local vars: inventory, weights]
	[Stack Base - Memory Address 0x0000]

Figure 2: LIFO Call Stack During Recursive QuickSort Execution

C++ Stack Demonstration Code:

```
stack<string> callStack;
callStack.push("main()"); // Frame 1 pushed
callStack.push("loadTruck()"); // Frame 2 pushed
```

```

callStack.push("quickSort()"); // Frame 3 pushed
callStack.push("partition()"); // Frame 4 pushed (TOP)

cout << "Current top: " << callStack.top(); // "partition()"
// Functions return in reverse order (LIFO):
while (!callStack.empty()) {
    cout << "POP: " << callStack.top(); // partition -> quickSort -> loadTruck -> main
    callStack.pop();
}

```

Such a LIFO mechanism is an absolute must for recursion in our QuickSort implementation. A stack frame is created for each recursive call of quickSort(). The one that is deepest (base case) is completed first. Average case maximum recursion depth for QuickSort is $O(\log N)$, e. g. 1 000, 000 items would yield about a 20 frames deep call stack. A stack overflow error is raised if the stack goes beyond its allocated memory, hence keeping an eye on recursion depth is vital in production systems.

P3: Formal Specification for a Software Stack (Imperative Definition)

A formal specification is a mathematical description of the features of an ADT such as its pre-conditions and post-conditions. It is entirely independent of any programming language. A stack with LIFO operations in a cargo loading scenario is described by this definition: the last item loaded on a truck is the first item unloaded at the delivery point.

In the imperative style, an operation is a change of state, e. g. stack in state S is transformed by operation Op into state S . Pre-conditions indicate the situation before the operation can be performed, while post-conditions indicate the situation after completion.

Operation Signature	Pre-conditions	Post-conditions
init(capacity: Integer) → Stack	capacity > 0	Returns Stack S where $S.size = 0$ and $S.maxSize = capacity$. Memory allocated for capacity elements.

push(S: Stack, item: T) → Stack	S.size < S.maxSize (Stack is not full) If S.size == S.maxSize, throws Stack Overflow Exception.	item is placed at S.top position. S.size = S.size + 1. All previous elements unchanged.
pop(S: Stack) → T	S.size > 0 (Stack is not empty) If S.size == 0, throws Stack Underflow Exception.	Returns element at S.top. Removes that element. S.size = S.size - 1.
peek(S: Stack) → T	S.size > 0 (Stack is not empty) If S.size == 0, throws Stack Underflow Exception.	Returns element at S.top without any modification. S.size unchanged.
isEmpty(S: Stack) → Boolean	None (always valid)	Returns TRUE if S.size == 0, FALSE otherwise. Stack state unchanged.
isFull(S: Stack) → Boolean	None (always valid)	Returns TRUE if S.size == S.maxSize, FALSE otherwise. Stack state unchanged.

Table 5: Formal Stack Specification Using Imperative Definition

The formal specification described here is not tied to any particular implementation: it might be implemented by using an array (with a fixed size and $O(1)$ access) or by a linked list (with a dynamic size and a pointer overhead). It only states WHAT the operations do, not HOW they do it this, is a key idea of abstraction in software engineering.

D1: Routing Analysis – Dijkstra's vs A* Algorithm

Swift-Load Logistics runs their operations by connecting warehouses, distribution centers, and delivery points through a network of roads that vary in distance and conditions. One of their priorities is to figure out the best delivery route that not only saves fuel but also cuts the delivery time and increases the number of deliveries that can be made in a day. We compare two well-known shortest-path algorithms that can be used on a distribution network graph.


Dijkstra's Algorithm:

Edsger Dijkstra's algorithm was introduced in 1959. It plans the shortest path from one source node to all other nodes in a weighted graph with non-negative edge weights. It keeps a priority queue of unvisited nodes sorted according to their tentative distances. At every step, it picks the unvisited node with the minimum tentative distance, marks it as visited, and changes the distances of its neighbors. Its evaluation function is $f(n) = g(n)$ where $g(n)$ means the real total cost from the source till node n . Since it explores all directions equally, it is therefore an uninformed or blind search algorithm.

A* Search Algorithm:

The A* (A-Star) algorithm, named after its three authors Hart, Nilsson and Raphael, dates back to 1968. It is a Dijkstra's algorithm plus a heuristic function $h(n)$ that predicts the remaining cost from node n to the goal. Its evaluation function is $f(n) = g(n) + h(n)$, where $g(n)$ is the actual cost from the source and $h(n)$ is the heuristic estimation. In route finding, the most commonly used heuristics are the Euclidean distance (i. e. straight line) and the Manhattan distance. When a heuristic is admissible (meaning it does not overestimate the actual cost), A* will not only be able to find the optimal shortest path but it will also explore less nodes as compared to Dijkstra, as it prioritizes nodes that seem to be on the way to the goal whereas Dijkstra explores all directions equally.

Figure 3: Swift-Load Regional Distribution Map



Following Dijkstra algorithm: Dijkstra is a blind uniform search. Starting at Main Warehouse (A), it first explores node B as it is the cheapest edge immediately available (10km). Then it goes to node C (15km). From node B it examines the path to D (total 10+20=30km). Eventually from node C it checks the path to D (total 15+10=25km). Apart from finding the optimal path (A -> C -> D), it also did a lot of extra calculations that led it to explore North Hub only because the very first step was the cheapest one.

Following A* algorithm: For simplicity let us assume that A* uses a straight-line Euclidean distance to the delivery point (D) as the heuristic. Then it will estimate the quality of paths by the function $f(n) = g(n) + h(n)$. Let the heuristic estimate from North Hub be $h(B) = 18\text{km}$, from South Hub be $h(C) = 8\text{km}$. Then A* will calculate $f(B) = 10 + 18 = 28$, and $f(C) = 15 + 8 = 23$. Since $23 < 28$, A* decides to go to South Hub (C). It actually sees

that over time South Hub is cheaper even though from A it is longer. It goes straight to the A -> C -> D path ignoring North Hub totally and the number of nodes explored also is A lot reduced.

Feature	Dijkstra's Algorithm	A* Algorithm
Evaluation Function	$f(n) = g(n)$	$f(n) = g(n) + h(n)$
Search Type	Uninformed (blind expansion)	Informed (heuristic-guided)
Time Complexity	$O(V^2)$ or $O((V+E) \log V)$	$O(E)$ best case with good heuristic
Space Complexity	$O(V)$	$O(V)$ for open/closed sets
Optimality	Always finds shortest path	Optimal if $h(n)$ is admissible
Nodes Explored	All reachable nodes (worst case)	Only nodes along promising paths
Best Use Case	Pre-computing ALL shortest paths	Real-time single-destination routing
Swift-Load Verdict	Use for overnight route pre-computation	Use for live GPS truck navigation

Table 6: Comprehensive Dijkstra vs A* Comparison for Delivery Routing

For the Swift-Load GPS navigation system, A* is the preferred choice as it significantly decreases the computation time by concentrating the search toward the delivery location. Nonetheless, Dijkstra's algorithm is still of great use in overnight batch processing when all the routes between warehouses are pre-computed and cached for quick access during the day.

P6: Effectiveness Assessment via Big O Notation

Asymptotic analysis via Big O notation offers a mathematical framework independent of hardware, which can be used to assess how good an algorithm is. Big O characterizes the highest limit of an algorithm's increasing rate of work as the size of input

N goes to infinity. It is a technique to neglect constant factors and lower order terms in order to expose the basic behavior of an algorithm as N increases. In this way, engineers are able to compare algorithms on an equal footing without considering the difference in the facilities on which their performances have been measured, such as a modern server versus a truck embedded computer.

The main idea behind Big O is that constant multipliers do not matter at large scales. A method of complexity $O(N)$ executed on an old computer will exceed the performance of an $O(N)$ algorithm on a supercomputer once N becomes very large. Here, N for Swift-Load can denote millions of packages being handled during holiday times, so this mathematical assurance is a very important factor for the system.

Algorithm	Big O Complexity	Operations for N=100	Operations for N=1,000,000
AVL Tree Search	$O(\log N)$	~7	~20
Linear Search	$O(N)$	100	1,000,000
Bubble Sort	$O(N^2)$	10,000	1,000,000,000,000 (1 trillion)
QuickSort	$O(N \log N)$	~664	~20,000,000
Dijkstra (priority queue)	$O((V+E) \log V)$	~1,400	~40,000,000

Table 7: Big O Scaling Comparison Across System Algorithms

P7: Two Methods for Measuring Algorithm Efficiency

Algorithm efficiency is evaluated along two complementary aspects: time complexity (duration of execution) and space complexity (amount of memory used). Both aspects matter a lot to the Swift-Load system, where the cargo processing needs to be very fast to allow real-time operations and memory-efficient to work on truck computers.

Method 1: Measuring Execution Time (Time Complexity)

We used the chrono library in our C++ source code to record time elapsed during algorithm execution in nanoseconds. Two timestamp measurements are taken by `high_resolution_clock` immediately before and after running the algorithm, and `duration_cast` is used to convert the elapsed time to nanoseconds. Besides, we have monitored the number of comparisons and swaps precisely done by each algorithm, which gives us an additional hardware-independent metric along with the wall-clock measurement.

```
auto start = high_resolution_clock::now();
bubbleSort(weights); // Algorithm under test
auto stop = high_resolution_clock::now();
auto duration = duration_cast<nanoseconds>(stop - start);
cout << "Execution Time: " << duration.count() << " nanoseconds";
// Results: Bubble Sort: 66 comparisons, 34 swaps
//          QuickSort:   34 comparisons, 23 swaps
```

Method 2: Memory Usage Analysis (Space Complexity)

Space complexity is gauged by assessing the memory size occupied by data structures with C++ `sizeof` operator and by considering the extra space taken up by algorithms. In our implementation, `sizeof(Goods)` results in around 80 bytes per object, while `sizeof(AVLNode)` is roughly 104 bytes because of the additional left pointer, right pointer, and height integer field. An array of 8 goods items is ~640 bytes of contiguous memory, while 8 AVL nodes need ~832 bytes plus pointer overhead. This exemplifies the fundamental space-time trade-off: the AVL tree consumes about 30% more memory per item yet provides $O(\log N)$ search performance as opposed to $O(N)$ for a linear array scan.

Even though Big O notation describes the theoretical efficiency of a program well, it ignores some important hardware factors, the most important one being the CPU Cache Locality. For example, arrays (like those used in our linear search) are laid out continuously in memory, which means they are very cache-friendly. The CPU, in fact, can load the next several array elements into the L1 cache at a high speed. But, trees (our AVL tree example, for instance) depend on pointers to nodes that are haphazardly allocated in the heap, leading to frequent and costly CPU cache misses. In fact, when dealing with very small data sets

(say $N < 50$ items), a linear search $O(N)$ through an array may regularly beat a tree search $O(\log N)$ solely due to hardware cache architecture differences. Deciding between them well calls for knowledge both of math and of silicon.

M5: Trade-off Interpretation in ADT Specification

A trade-off, in software engineering, is when improvement of one feature of the system results in the degradation of another feature. In the realm of data structures, the most basic and major trade-off is between time efficiency (speed of operations) and space efficiency (memory consumption). All design choices in the Swift-Load system are time-space trade-off decisions aimed at the highest operational efficiency of the system.

Concrete Example: Choosing Between Hash Table, AVL Tree, and Linked List

We considered 3 candidate data structures before selecting the final one for the warehouse inventory system. They were located on different positions on the time-space trade-off spectrum:

Data Structure	Search Time	Memory Per Item	Trade-off Analysis
Hash Table	$O(1)$ average	High (~200 bytes: data + bucket + hash)	Fastest possible search but wastes significant memory on empty buckets and hash storage. Load factor management adds complexity.
AVL Tree (chosen)	$O(\log N)$	Moderate (~104 bytes: data + 2 pointers + height)	Excellent search speed with reasonable memory overhead. Self-balancing guarantees worst-case $O(\log N)$. Ordered traversal is free.
Linked List	$O(N)$	Low (~88 bytes: data + 1 pointer)	Most memory-efficient per node but search time is linear. Unacceptable for

			real-time package tracking with thousands of items.
--	--	--	---

Table 8: Space-Time Trade-off Comparison for Inventory Management

For Swift-Load Logistics, an AVL Tree has been selected as the best balance. Although the Hash Table offers $O(1)$ search time that looks great on paper, its memory waste due to empty buckets and the complicated nature of hash collision resolution make it unsuitable for the memory-limited embedded computers installed in the delivery trucks. On the other hand, while the Linked List's very low memory usage is highly advantageous, its $O(N)$ search time would cause multi-second delays when locating packages in a warehouse with tens of thousands of items, a clear violation of the real-time performance requirements. Thus, the AVL Tree's $O(\log N)$ search with moderate memory overhead perfectly fits the criteria of an ideal middle ground, delivering sub-millisecond search times, yet comfortably residing within the computer's RAM available on the truck.

Task 3: Implementation and Technical Critique

The last assignment is aimed at hand-developing a highly difficult self-balancing tree data structure, thoroughly testing it for robustness, and performing an academic peer review of its performance through formal Big O complexity analysis. The coding part illustrates the transition of computer science theory into software which effectively handles logistics management problems.

P4 & M4: Complex ADT Implementation – AVL Tree for Warehouse Inventory

The AVL Tree (after its creators Georgy Adelson-Velsky and Evgenii Landis, 1962) is a self-balancing Binary Search Tree (BST) which enforces a particular balance rule very

strictly: at each node of the tree, the difference in heights of the left and right subtrees (balance factor) should be, in absolute value, no more than 1. This condition ensures that the height of the tree is never more than $1.44 * \log_2(N)$, and thus insertion, deletion, and search operations can be performed in $O(\log N)$ time in the worst case. On the other hand, a normal BST may deteriorate to $O(N)$ if elements are inserted in sorted order, hence producing a linear chain (which is just like a linked list).

For warehouse inventory at Swift-Load, AVL's assurances are essential. As we manage thousands of parcels, we simply cannot risk a scenario where a degenerate BST turns our $O(\log N)$ searches into $O(N)$ linear scans. Thanks to its automatic rebalancing, the AVL tree completely removes this possibility.

How Self-Balancing Works – The Four Rotation Cases:


The insertion operation causes the tree to update the balance factor ($BF = \text{height}(\text{left}) - \text{height}(\text{right})$) for all the nodes that are on the path from the newly inserted node back to the root. In case the balance factor of a node falls out of the interval $[-1, +1]$, rotations are carried out to bring it back to balance. In fact, there are only four situations:

Imbalance Case	Detection Condition	Rotation Applied	When It Occurs
Left-Left (LL)	$BF > 1$ AND new key < left child key	Single Right Rotation at unbalanced node	Consecutive insertions into the left subtree of the left child
Right-Right (RR)	$BF < -1$ AND new key > right child key	Single Left Rotation at unbalanced node	Consecutive insertions into the right subtree of the right child
Left-Right (LR)	$BF > 1$ AND new key > left child key	Left Rotate left child, then Right Rotate node	Insertion into the right subtree of the left child
Right-Left (RL)	$BF < -1$ AND new key < right child key	Right Rotate right child, then Left Rotate node	Insertion into the left subtree of the right child

Table 9: AVL Tree Four Rotation Cases

To demonstrate exactly how the implementation solves tree degradation, consider the step-by-step visual transformation during a Left-Left Imbalance. Inserting items in reverse sorted order (3, then 2, then 1) would normally create a degenerate linear chain. The AVL tree detects this and fixes it via a Right Rotation:

AVL Tree 'Left-Left Imbalance'




Rotation Implementation (Right Rotation Example):

```
AVLNode* rightRotate(AVLNode* y) {
    AVLNode* x = y->left;      // x becomes new root of subtree
    AVLNode* T2 = x->right;     // Save x's right subtree
    x->right = y;               // y becomes right child of x
    y->left = T2;               // T2 becomes left child of y
    // Recalculate heights (bottom-up)
    y->height = max(height(y->left), height(y->right)) + 1;
    x->height = max(height(x->left), height(x->right)) + 1;
```

```
    return x; // x is the new root of this subtree
}
```

AVL Tree Structure Produced by Our Program (8 Items):



The tree output above depicts 8 warehouse items stored in the completely balanced manner with the maximum height of the tree is 4. Each node has the balance factor greater than or equal to -1 and less than or equal to +1, that is why the AVL invariant has been kept. The root node 'Desk' was actually not the first item inserted; instead, it has been automatically promoted to root through rotations to minimize the overall tree height. The in-order traversal arrangement of the elements in the alphabetical order (Apples Bananas

Chair, Desk Laptop Monitor Printer Tablet) proves that the BST ordering property has been preserved along with balancing.

Our current solution is capable of providing extremely fast data retrieval to our logistics system. A tree with the height of 4 for only 8 items means the worst-case search is done with at most 4 comparisons. If there are 1 000 000 items in the warehouse, the maximum height of the tree will be 29, so any package can be located in fewer than 29 comparisons - an almost instantaneous operation even on very low power hardware.

P5: Robustness and Testing Report

Robustness refers to software's capability to work correctly in the presence of errors, unexpected or even harmful inputs without crashing, giving wrong results or becoming in a state which is undefined. In a logistics system where data is coming from different sources such as barcode scanners, manual entry terminals, and automated warehouse systems, invalid data is not just a matter of theory; it is the operational reality of each day. Our solution is based on defense-in-depth approach: several layers of error handling are associated with the system.

Error Handling Mechanisms Implemented:

1. **Constructor Validation:** During Goods object construction, the weight attribute is checked for being non-negative, and the throwing of `std::invalid_argument` is done if the condition is violated. This means that it is impossible to generate corrupted objects.
2. **Setter Validation:** The `setWeight()` and `setName()` functions check the inputs when before altering the internal state, they deny negative weights and empty names with meaningful exceptions.
3. **AVL Insertion Guard:** The `insertNode()` method verifies that the `CString` name is not empty before it does BST insertion, thereby the tree's key structure is not corrupted by the insertion of an unnamed item.

4. Try-Catch Wrappers: Recursive insertNode() is called by the public insert() method inside a try-catch block, which guarantees that exceptions have been caught and logged rather than the whole application being crashed by deep recursion.

```
// Layer 1: Constructor validation
Goods(string name, string type, double weight) {
    if (name.empty())
        throw invalid_argument("Name cannot be empty");
    if (weight < 0)
        throw invalid_argument("Weight cannot be negative");
}
// Layer 2: AVL insertion guard
AVLNode* insertNode(AVLNode* node, Goods item) {
    if (item.getName().empty())
        throw invalid_argument("Cannot insert item with empty name.");
    // ... BST insertion and rotation logic ...
}
// Layer 3: Public method with try-catch
void insert(Goods item) {
    try {
        root = insertNode(root, item);
    } catch (const exception& e) {
        cout << "[ERROR] " << e.what() << endl;
    }
}
```

Test Results Report:

Test ID	Test Description & Input	Expected Behavior	Result
T1	Create Goods with weight = -5.0 Goods("Bad", "X", -5.0)	Exception: "Weight cannot be negative"	PASS ✓
T2	Insert item with empty name insert(Goods("", "Misc", 10))	Error logged, tree unchanged	PASS ✓
T3	Set weight to -100 via setter g.setWeight(-100)	Exception: "Weight cannot be negative"	PASS ✓
T4	Set name to empty string g.setName("")	Exception: "Name cannot be empty"	PASS ✓

T5	Search for non-existent item search("Keyboard")	nullptr returned, no crash	PASS ✓ (4 comparisons)
T6	Search for existing item search("Monitor")	Valid Goods pointer returned	PASS ✓ (2 comparisons)
T7	Insert 8 valid items sequentially	Tree balanced, height = 4, all BF in [-1,1]	PASS ✓
T8	In-order traversal after insertions	Items displayed in alphabetical order	PASS ✓
T9	Insert Duplicate Key\ninsert("Monitor")	Exception: 'Key already exists'	PASS: The AVL tree prevents dictionary corruption by rejecting identical primary keys with a specific invalid_argument exception.

Table 10: Complete Robustness Test Results (8/8 Tests Passed)

D3: Critical Evaluation – AVL Tree vs Linear Search Using Big O

This section aims to provide a very detailed formal complexity analysis comparing the implemented AVL tree search algorithm against a simple linear search on an unsorted array. The comparison is done by both theoretical Big O analysis and empirical measurements from our program output to show the exponential superiority of balanced tree structures at scale.

Empirical Results from Program Execution (8 items):

Search Key	AVL Tree Comparisons	Linear Search Comparisons	AVL Improvement Factor
Monitor	2	4	2.0x faster
Tablet	3	6	2.0x faster

Bananas	2	8	4.0x faster
Keyboard (missing)	4	8	2.0x faster
Average	2.75	6.5	2.4x faster

Table 11: Empirical Search Performance Comparison (N = 8 items)

Although the empirical advantage appears minor when dealing with only 8 items, the theoretical difference becomes so huge when N is scaled to production levels:

Theoretical Scaling Analysis:

Dataset Size (N)	AVL Search: $O(\log N)$	Linear Search: $O(N)$	AVL Advantage
8 (our test)	~3 comparisons	~4 average	1.3x
1,000	~10 comparisons	~500 average	50x
100,000	~17 comparisons	~50,000 average	2,941x
1,000,000	~20 comparisons	~500,000 average	25,000x
1,000,000,000	~30 comparisons	~500,000,000 average	16,666,667x

Table 12: Theoretical Scaling – AVL $O(\log N)$ vs Linear $O(N)$

Formal Big O Comparison:

Complexity Metric	AVL Tree	Linear Search (Array)
Best Case Search	$O(1)$ – root is target	$O(1)$ – first element is target
Average Case Search	$O(\log N)$	$O(N/2) = O(N)$
Worst Case Search	$O(\log N)$ – guaranteed by balance	$O(N)$ – target is last or missing
Insertion Cost	$O(\log N)$ with rotations	$O(1)$ append to end
Space per Element	~104 bytes (data + pointers + height)	~80 bytes (data only)
Ordered Traversal	$O(N)$ in-order – free alphabetical listing	$O(N \log N)$ – must sort first

Table 13: Formal Big O Complexity Comparison

Key Takeaway: An AVL Tree is dramatically better in any scenario where search operations are much more than insertions, which precisely aligns with the case of warehouse inventory management. Items (packages) in a warehouse are scanned (searched) many times over the lifecycle but they are inserted only once. The slight insertion overhead of $O(\log N)$ and 30% extra memory per node are negligible more than compensated by the huge change in search efficiency going from $O(N)$ to $O(\log N)$. In a real warehouse environment having 1 000 000 items, the difference will be 20 comparisons versus 500,000 comparisons for each query - a 25,000x improvement that changes multi-second wait time to sub-microsecond response.

D4: Three Benefits of Implementation-Independent Data Structures

Implementation independence is one of the key software engineering principles which essentially separates the behavioral specification of a data structure at an abstract level (WHAT it does) from its concrete realization at the code level (HOW it does it). One way to achieve this separation is through interfaces that not only define method signatures but also specify parameter types, return types, and behavioral contracts without disclosure of internal implementation details. For the Swift-Load Logistics platform, following this principle yields three main benefits:

Benefit 1: Cross-Platform Portability

As the Goods ADT and AVL Tree only expose three abstract interface methods (insert search display) and do not reveal their C++ implementation details, it is possible to port the entire system architecture to different programming languages and platforms without redesigning the core logic. Tables 1 and 5 of the formal specification not only describe the system's inputs and outputs but also specify its operational behavior and provide detailed sequences of operations. These tables are language-agnostic blueprints that any competent developer can implement in Java Python TypeScript, or Rust. Consequently, Swift-Load can be diversified from a desktop C++ application to a web-based

dashboard(TypeScript/React), mobile tracking app(Swift/Kotlin), or cloud-native microservice(Go/Rust) without the need for a complete architectural rewrite. Hence, the investment in formal specification is rewarded each time the system is ported to a new platform.

Benefit 2: Internal Refactoring Without Breaking Changes

A Internal Refactoring Without Breaking Changes The AVL Tree implementation can be replaced entirely by a few lines of code by using a Red-Black Tree, B-Tree, or Skip List for different workload characteristics by a senior developer. For instance, if the analysis shows that the warehouse workload is write-heavy (lots of insertions, fewer searches), the Red-Black Tree, which has cheaper rotations, could be a better choice. Since external code interacts only with public interface methods `insert()` and `search()`, which are the same, no matter the underlying tree type, no changes are required in other modules. The truck-loading module, billing system, delivery dashboard, and reporting engine all continue working without any changes and without recompilation, significantly lowering regression risk and the cost of testing.

Benefit 3: Parallel Team Development and Reduced Integration Risk

One of the advantages of implementation independence is that it allows that multiple teams can genuinely develop concurrently. For example, the frontend team, while designing the warehouse dashboard UI, could assume that the `search()` function takes a string key and returns a Goods pointer (or null if not found). At the same time, the backend team could be deeply engaged in implementing the complex AVL tree rotation logic without any interference. Both teams rely on the same formal interface contract that is documented in the specification tables. On the day of integration, these components just plug in nicely with each other because both sides have adhered to the agreed-upon interface. According to some industry reports, such a separation of concerns can in fact cut the total development time by 30-40% in case of a team of 5+ developers since it removes the blocking dependencies where one team has to wait for the other's implementation, which results in the latter being able to start their work only after the completion of the former's.

Code-Level Realization in C++

While implementation independence is a vital theoretical concept, it must be grounded in an executable programming language to fulfill the distinction criteria. In C++, implementation independence is formally achieved using Abstract Base Classes (ABCs) containing Pure Virtual Functions. By defining an interface such as ``virtual void insert(Goods item) = 0;`` and ``virtual Goods* search(string key) = 0;``, the C++ compiler strictly enforces the contract. Any class inheriting from this interface must provide the concrete implementation, ensuring that the backend tree structure (whether AVL, Red-Black, or B-Tree) remains entirely decoupled from the frontend components calling it. Stating this exact language feature proves that the architectural abstraction is successfully mapped to executable code.

References

Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2009). Introduction to algorithms (3rd ed.). MIT press.

Sedgewick, R., & Wayne, K. (2011). Algorithms (4th ed.). Addison-Wesley Professional.

Stroustrup, B. (2013). The C++ programming language (4th ed.). Addison-Wesley.

Hart, P. E., Nilsson, N. J., & Raphael, B. (1968). A formal basis for the heuristic determination of minimum cost paths. IEEE transactions on Systems Science and Cybernetics, 4(2), 100-107.

Drepper, U. (2007). What every programmer should know about memory. Red Hat, Inc.

Appendix

GitHub Repository Link for C++ Source Code (main.cpp):

<https://github.com/xshadowhunterx007-prog/DSAAssignment/>

The complete C++ source code (main.cpp) is available at the GitHub link above. The program implements all data structures and algorithms discussed in this report, including the Goods ADT, FIFO Queue, Stack demonstration, Bubble Sort, QuickSort, AVL Tree with self-balancing rotations, linear search for comparison, robustness testing, and execution time/memory measurement. The terminal output (output.txt) contains all test results and performance data referenced throughout this document.